COMP15111 notes

Todd Davies

October 25, 2019

Note, extra space has been allocated for the right hand margin to allow for more extensive margin notes. Also, it gives you space to make your own annotations.

These notes are based off the material from the COMP15111 course run by Dr. Javier Navaridias-Palma and Dr. Richard Neville. They are released under a Creative Commons licence, please submit issues and pull requests at `https://github.com/Todd-Davies/first-year-notes`.

Contributors

- Chris Williamson (Github: ChrisWilliamson123)
- Sami Alabed (Github: SamiAlabed)

Contents

1   Introduction

1.1   A Computational Model

The simplest, earliest, commonest, most important computational model is the **Von-Neumann Imperative Procedural Computer Model**

According to this model, a computer can:

1. Store information

2. Manipulate the stored information

3. Make decisions depending on the stored information

1.2   Simple View Of A Computer

The simplest model of a computer can be represented as:

$$Memory \Leftrightarrow Bus \Leftrightarrow Processor$$

1.2.1   Memory

Memory is a set of locations which can hold information, such as numbers, programs or many other types of data. Each memory location has a unique numerical address, and there are typically thousands of millions of different locations. There are various ways of depicting memory; a common one is a 'hex dump' that often looks something like this:

Address	Hex values	ASCII
00000000	48 65 6C 6C 6F 0A	Hello.

Each item that is in the memory has a unique address.

Run the command `hexdump` to generate hexdumps.

1.2.2   Bus

A bus is a bidirectional communication path. It is able to transmit addresses and numbers between components inside the computer.

1.2.3   Processor

The processor obeys a sequence of instructions, commonly referred to as a program. Historically the processor was often referred to as a CPU, however, this is inappropriate nowadays since typical processors consist of several processing cores.

1.3   Three-address instructions

Every kind of processor has a different set of instructions, real world examples include: Pentium, ARM and others

Each three-address instruction:

1. Copies the values from any two memory locations and sends them to the processor (source operands)

2. Copies some operation e.g. adds the copied numbers together

3. Copies the result back from the processor into a third memory location (destination operand)

4

For example, if we wanted to convert the Java code `sum = a + b;` into a three-address instruction we would:

1. Identify the two *source operands*: *a* holds 2, *b* holds 3

2. Perform the *operation*: $2 + 3 = 5$

3. Let the variable *sum* equal the answer 5. This is the *destination operand*

1.3.1 Three address example

Question: Convert the Java code *product* = *c* * *d;* into the three- address style and draw a two box view of it.

First we need to re-write the Java code in the three-address style:

$$product \leftarrow c * d$$

Now we can draw the box view of it (figure 1).

Figure 1: An example of the two box model

1.3.2 Memory bottleneck

Most processors can process instructions faster than they can be fed by memory. Each instruction in the three-address cycle requires four memory cycles:

1. Fetch the instruction

2. Read the first operand

3. Read the second operand

4. Write the result to memory

Each of these memory cycles could take hundreds of processor clock cycles to complete, and so in this time the processor would be doing nothing. However, most modern processors employ a *cache* to temporarily store commonly accessed memory locations, and so avoid some of the memory cycles.

1.4 Registers

Registers are very small amounts of storage build into a processor. Since they are inside the processor data doesn't need to be transferred over the bus, and so they are very fast. Registers are used instead of the main memory which speeds up program execution.

Each register can only hold one value and each processor will only generally have a few dozen registers (e.g. ARM has sixteen).

1.5 Instruction Styles

1.5.1 One address

The one address style can only use up to one memory location in each instruction, all other operands must be registers. An example may be:

$$R1 \leftarrow R0 + memory\ location$$

1.5.2 Load-store

The load-store style cannot perform operations on memory locations at all. Instead, values from memory must be loaded into a registers before the operation takes place and then the operation can be performed on the registers. Following the operation, the result must be stored back into memory again.

$$R1 \leftarrow memory\ location$$
$$R1 \leftarrow R0 + R1$$
$$memory\ location \leftarrow R1$$

This means that we need extra instructions to do stuff with memory locations:

1. **Load** the value from memory into a register before the operation.

2. **Store** the value in the register back to memory after the operation.

For example, the Java code *Sum = a + b + c;* would be run as:

R1	←	a	(i.e. load from a)
R2	←	b	(i.e. load from b)
R3	←	R1 + R2	(i.e. a+b)
R4	←	c	(i.e. load from c)
R5	←	R3 + R4	(i.e. (a+b)+c)
Sum	←	R5	(i.e. store to sum)

You can see that the load-store style favours lots of very simple, very fast instructions.

1.6 An introduction to bases

Conventionally, we count using base 10. Base 10 includes, you guessed it, ten different symbols from 0 through to 9.

Sometimes however, it is convenient to count using different bases. Popular bases include:

Base n	Member symbols	Name
$n = 2$	$\mathbb{Z}_2 = \{0, 1\}$	Binary
$n = 8$	$\mathbb{Z}_8 = \{0, 1, 2, 3, 4, 5, 6, 7\}$	Octal
$n = 10$	$\mathbb{Z}_{10} = \{0, 1, \ldots, 7\}$	Decimal
$n = 16$	$\mathbb{Z}_{16} = \{0 - 9, A - F\}$	Hexadecimal

1.6.1 How to read numbers in any given base

The formula for reading a number in a given base is as follows:

$$\sum_{i=0}^{k} a_i b^i$$

Where the number you're trying to read takes the form $a_k, a_{k-1}, \ldots, a_2, a_1, a_0$, b is the base you're using and i is the count from right to left of the digits in the number.

Example 1 Lets apply the formula to the base 10 number 27385:

$$\begin{aligned}
27385 &= (5 \times 10^0) + (8 \times 10^1) + (3 \times 10^2) + (7 \times 10^3) + (2 \times 10^4) \\
&= (5 \times 1) + (8 \times 10) + (3 \times 100) + (7 \times 1000) + (2 \times 10000) \\
&= 5 + 80 + 300 + 7000 + 20000 \\
&= 27385
\end{aligned}$$

Example 2 Lets apply the formula to the base 16 number F00BA4:

$$\begin{aligned}
F00BA4 &= (4 \times 16^0) + (A \times 16^1) + (B \times 16^2) + (0 \times 16^3) + (0 \times 16^4) + (F \times 16^5) \\
&= (4 \times 16^0) + (10 \times 16) + (11 \times 256) + (0 \times 4096) + (0 \times 65536) + (15 \times 1048576) \\
&= 4 + 160 + 2816 + 0 + 0 + 15728640 \\
&= 15731620
\end{aligned}$$

1.6.2 Changing from base 10 to base n

In order to change into base n from base 10, we just repeatedly divide by n and use the remainder as the value for base n. Here are a few examples:

Example 1 Convert 893 into base 2.

$$
\begin{array}{rcll}
893 \div 2 & = & 446 & r1 \\
446 \div 2 & = & 223 & r0 \\
223 \div 2 & = & 111 & r1 \\
111 \div 2 & = & 55 & r1 \\
55 \div 2 & = & 27 & r1 \\
27 \div 2 & = & 13 & r1 \\
13 \div 2 & = & 6 & r1 \\
6 \div 2 & = & 3 & r0 \\
3 \div 2 & = & 1 & r1 \\
1 \div 2 & = & 0 & r1 \\
\end{array}
$$

Reading up from the bottom, we can see that the binary (base 2) representation is 1101111101.

Example 2 Convert 893 into base 9.

$$
\begin{array}{rcll}
893 \div 9 & = & 99 & r2 \\
99 \div 9 & = & 11 & r0 \\
11 \div 9 & = & 1 & r2 \\
1 \div 9 & = & 0 & r1 \\
\end{array}
$$

Reading up from the bottom, we can see that the nonal (base 9) representation is 1202.

Example 2 Convert 893 into base 16.

$$
\begin{array}{rcll}
893 \div 16 & = & 55 & r13 \\
55 \div 16 & = & 3 & r7 \\
3 \div 16 & = & 0 & r3 \\
\end{array}
$$

Reading up from the bottom, we can see that the hexadecimal (base 16) representation is $3, 7, 13$ or $37D$.

2   ARM

Computers obey programs which are sequences of instructions. Instructions are coded as values in memory. The sequences are held in memory adjacent memory locations. Values in memory can be interpreted as you please, from numbers to text, images or anything really!

Any given set of binary digits can be read as a decimal number, but not always as text, so values in memory are often represented as numbers for convenience.

2.1   Assembly Language

Assembly language is a means of representing machine instructions in a human readable form.

Each type of processor has its own assembly language (since each language is specific to a particular architecture) but each instruction typically has a lot in common:

- A mnemonic, that specifies the type of operation
- A destination, such as a register or memory location
- And one or more sources that may be registers or memory locations.
- Possibly with a comment too which will help programmers understand what's happening and aren't interpreted by the assembler.

When a program has been written in assembler, it must be *assembled* by an *assembler* to run it.

2.2 ARM instructions

ARM has many instructions but we only need three categories:

- Memory operations that move data between the memory and the registers.
- Processing operations that perform calculations using value already in registers.
- Control flow instructions are used to make decisions, repeat operations etc.

2.3 Transferring data between registers and memory

Memory operations load a value into a register from an address in memory or store the value of a register to a memory address.

For example, a into register 1 ($R1 \leftarrow a$) we would write: `LDR R1, a`

Or to store the value in register 5 into sum ($sum \leftarrow R5$): `STR R5, sum`

In these examples, a and sum are aliases for the addresses of memory locations.

2.4 ARM processing instructions

ARM has many different instructions to perform operations such as addition, subtraction and multiplication.

The syntax for such operations is usually:

`[operand] [destination register] [register 1] [register 2]`

For example, to add two numbers together, we might write:

`ADD R2, R0, R1`

This will add the value of R0 to the value of R1 and store it in R2.

2.5 ARM control instructions

The most common control instruction is the branch. Similar to `GOTO` in other languages, a branch will change the PC register (see section 2.6) to another value so the order of execution of the program is changed.

Branches can be made to be conditional by appending a conditional operator on to the command.

The syntax is something like:

`B[conditional operator] [branch name]`

Some examples of different conditional operators are:

Command	Function
B	Branches to a different location in the code.
BNE	Branches, but only if the previous condition was false.
BEQ	Branches, but only if the previous condition was true.

2.6 Stored programs and the Program Counter

A computer can make decisions, and choose which instructions to obey next depending upon the results of those decisions. A **Program Counter** (PC) register is used to hold the memory address of the next instruction to be executed. ARM uses register 15 as its PC.

2.7 Fetch-Execute Cycle

The processor must first fetch instructions from memory before it can execute them. This is called the fetch-execute cycle, and it involves:

1. **Fetch**: copy the instruction, pointed to by the PC, from memory and set PC to point to the next instruction

2. **Execute**: obey the instruction (exactly as before)

3. Repeat.

In ARM, the PC starts with a values of `0x00000000` when the program is initially run. On each cycle of the Fetch-Execute cycle, the PC is incremented by 4, since instructions each occupy 4 memory locations.

2.8 Decision Making

In order to make decisions, the computer mustn't just execute instructions one after the other in a linear manner. Instead, branches must be used to change the sequence of instructions to be executed.

In order to perform a conditional branch, we must first perform a compare command to perform the comparison before we do the branch.

2.8.1 An example

If we wanted to do a $1 discount on a shopping list if the price was over $20, we would do the following:

```
1                      LDR     R0, total       ; Load the total price
2                                              ; into R0
3
4                      CMP     R0, #20         ; Compare R0 and 20
5                                              ; (the literal)
6                      BLT     nodiscount      ; If the price is too low,
7                                              ; then don't discount
8                      SUB     R0, #1          ; Deduct £1
9                      STR     R0, total       ; Store the result back
10                                             ; into memory
11
12   nodiscount        SVC     2               ; Finish
13
14   total             DEFW    25              ; Lets say the total is \$25
```

2.9 Allocating memory

The `DEFW` (define word) operation puts a value in memory before the program is run. Any define operation is executed before the program is run.

The actual memory location that is used to store the value isn't known to the running program, however, an *alias* is attached to the memory location by the programmer and the memory location can be referenced through that.

The syntax for the `DEFW` command is as follows:

`myage DEFW 18`

Where `myage` is the alias and `18` is the value.

`DEFW` can also be used to define a number of words:

`squares DEFW 0, 1, 4, 9, 16, 25`

The label (`squares`) is associated with the lowest address (i.e. 0)

`DEFB` stores a single byte in memory. It is useful for strings such as "hello":

`hi DEFB "hello"`

`DEFS` sets a block of bytes to a set value:

`reserved_space DEFS 10, 5`

The above will set 10 bytes to the value '5'.

3 Storing values

There are many ways to store data. For example, we could store what lights are one in a traffic light in many different ways. First, we must decide how many different states the traffic light can be in:

`Red, Red Amber, Green, Amber`

You can see that we have four states. This could be represented in binary as two bits:

00	Red
01	Red Amber
10	Green
11	Amber

We could also store the states as their binary representations of their names:

R	01010010
RA	0101001001000001
G	01000111
A	01000001

You can see though, that this isn't as efficient as storing just two binary digits.

4 ARM assembly programming

4.1 Different types of values

ARM has the capacity to work with many different types and sizes of values. Each type has a different use case. The main ones are described below:

Name	length	Use
Byte	8 bits	Used for characters
Word	32 bits	Used for integers, addresses and instructions

There are other types too (such as the halfword and doubleword) but they aren't needed for this module.

ARM processors require that memory locations are aligned. This means that values stored in memory start at specific places. For example, a word address must be a multiple of four. e This means that after a DEFW statement, the ALIGN command must be called (See 4.10.2 for more on the ALIGN command.).

4.2 Loading and storing values in memory

The commands LDR and STR are used to move values between memory and registers. The commands are detailed in full below:

Command	Function
STR	Copies the whole (32 bit) register into memory.
LDR	Loads a 32 bit word from memory into a register.
STRB	Stores a single 8 bit byte into memory from a register.
LDRB	Loads a byte from memory into a register. The upper 24 bits of the register are zeroed.

4.3 Endianness

Endianness defines how bytes should be ordered in multi-byte values. There are two types of endianness: **little endian** and **big endian**.

In the word 0x12345678 there are four bytes:

- 0x12

- 0x34

- 0x56

- 0x78

In little endian, the first byte would be 0x78 since bytes are read from right to left.

In big endian, the first byte would be 0x12 since bytes are read from left to right.

Although it does not matter which endianness is chosen, it must be kept consistent across all read and writes of a location in memory. Otherwise, the read value will differ from the written value e.g. 0x12345678 may be written as 0x12 0x34 0x56 0x78 (big endian) and then it may later be interpreted as 0x78563412 (little endian).

In this course, little endian is used, though ARM can use either.

4.4 Addressing memory

ARM uses 32 bit addresses, so there are 2^{32} different bytes that can be addressed in memory (or $\frac{2^{32}}{4}$ different words). However, there is no guarantee that the system on which the program is running will have that much memory available.

4.5 Instruction encoding

Each ARM instruction is encoded into a four byte word. The exact meaning of each of the bits varies per instruction.

For example, in the branch instruction, the first four bits specify the condition, the second four bits represent the actual operation to perform (i.e. branch) and the remaining twenty four bits define the memory location of the next instruction to branch to.

However, this presents a problem. We only have twenty four bits with which to define the next location to branch to, which allows us to define 2^{24} different locations. However, there are 2^{32} possible addresses that we could use!

This problem is overcome by treating the 24 bits as an offset to the address of the current instruction. This works since most of the time, the address that is being branched to is fairly close to the current instruction.

In order to be able to branch to addresses before and after the current instruction, we must use two's complement to allow signed integers to be used to specify the offset. This means we can branch to any instruction at an address $\pm 2^{23}$ from the current instruction.

4.6   Literals

ARM is able to encode literal values into instructions. This saves time having to access registers or memory in order to perform operations such as arithmetic.

An example is to increment a register:

```
ADD R1, R1, #1
```

However, ARM only assigns up to 12 bits for a literal value, so we can only have 2^{12} values. However, ARM employs a strange method of encoding these values so that more useful values are available (for example, #512 is allowed, but #257 isn't).

4.6.1   Negative literals

Technically, ARM doesn't support negative literals, however, the assembler will usually be able to find a way to implement them. Some examples are given below:

$$
\begin{array}{lcl}
\text{ADD R1, \#-1} & \rightarrow & \text{SUB R1, \#1} \\
\text{CMP R2, \#-2} & \rightarrow & \text{CMN R2, \#2} \\
\text{MOV R3, \#-3} & \rightarrow & \text{MVN R3, \#3}
\end{array}
$$

CMN is compare negative.
MVN is move not.

4.7   Supervisor calls

Supervisor calls are functions implemented by the operating system, not ARM itself. The parameter of an SVC call defines it's exact operation.

In this module, the SVC call does the following for each parameter:

SVC 0	Output a character
SVC 1	Input a character
SVC 2	Stop execution
SVC 3	Output a string
SVC 4	Output an integer

SWI is the old (yet still occasionally used) command for SVC - they do the same thing. SWI stands for *SoftWare Interrupt*, and SVC stands for *SuperVisor Call*

4.8   Pseudo instructions

The ARM assembler provides some instructions that are translated into sequences of more complicated instructions at the time of assembly for our convenience.

One such instruction is loading a literal into a register. This is done using the `LDR` command as usual, however a literal is used with the '=' character instead of a '#'. E.g. to load the value 100 into register one, we do:

```
LDR R1,=100
```

However, this is a pseudo instruction and will be converted by the assembler to:

```
MOV R1,#100
```

However, if the number is very large, it becomes:

```
constant DEFB 100
LDR R1, constant
```

4.9   Loading an address into a register

The `ADR` command loads an address into a register, for example:

```
1    constant        DEFB    100
2                    ADR     R1, constant
```

It will load the memory address of `constant` relative to the `PC` into `R1`. When the code is assembled, it is turned into two instructions (so `ADR` is a pseudo instruction). They are:

```
1            ADD R1, R1, PC
2            LDR PC, R1
```

4.10   Directives

Directives are evaluated at the time of assembly.

4.10.1   DEF commands

The `DEF{W,B,S}` command reserves an amount of memory dependent on the operation used (see the table) and puts an initial value in it.

Command	Function
DEFW num	Reserves a *word* of memory and puts the initial value `num` in it.
DEFB value	Reserves *byte*(s) of memory and puts the initial value `value` in it. Note that the value can be a string literal, in which case the number of bytes reserved will be equal to the length of the string.
DEFS size, fill	Reserves a *block* of memory of `size` bytes and initialises them with the value `fill`.

4.10.2   Align

The align command leaves as many blank bytes as needed so that the next item in memory will start at a word boundary (a multiple of 4).

4.10.3   Entry

Sets the PC at the start of the program (i.e. where the program should start from)

4.10.4   EQU

Allows you to name a literal, which can go a long way to making the code more maintainable. We could define the literal 18 as *drinking_age*:

```
1    drinking_age    EQU     #18
2                    ; Check the person is over 18
3                    CMP     R1, #drinking_age
4                    BLT     too_young
```

5   Arithmetic

5.1   Making good use of registers

Registers are a precious resource when programming on ARM. When writing software to evaluate expressions, it's often tempting to load all the variables into registers first, and then perform the arithmetic in separate registers like so:

```
1                ; a = b + c + d
2                LDR     R1, b
3                LDR     R2, c
4                LDR     R3, d
5                ADD     R4, R1, R2
6                ADD     R5, R4, R3
7                STR     R5, a
```

However, this is pointless - the values in the registers `R4` and `R5` aren't going to be needed again, so we may as well do:

```
1    ; a = b + c + d
2    LDR     R1, b
3    LDR     R2, c
4    LDR     R3, d
5    ADD     R1, R1, R2
6    ADD     R1, R1, R3
7    STR     R1, a
```

But, we can optimise even further here. Instead of loading all the variables into registers before we do arithmetic, we can save a register and load only the ones we need before each `ADD` instruction:

```
1    ; a = b + c + d
2    LDR     R1, b
3    LDR     R2, c
4    ADD     R1, R1, R2
5    LDR     R2, d
6    ADD     R1, R1, R2
7    STR     R1, a
```

Sometimes, it's useful to re-order an arithmetic expression so it can be implemented using less registers. This can often be achieved by increasing the nesting of brackets in an expression such as this:

$$(b - e) + (c * d) = ((c * d) + b - e)$$

Though these expressions are both equal, the right hand side will use a register less in ARM code, since each instruction can be executed sequentially, however the left hand side requires two expressions to be evaluated (and stored in a total of three registers) and then both expressions added together.

5.2 Using literals in expressions

If there is a literal in the expression you want to evaluate, then it's possible to use a literal instead of a register. For example, these two programs will end up with the same answer in `R0` but the one on the right uses less registers:

```
1          ; (a + 5) * b        1          ; (a + 5) * b
2          LDR    R0, five       2          LDR    R0, a
3          LDR    R1, a          3          ADD    R0, R0, #5
4          ADD    R0, R0, R1     4          LDR    R1, b
5          LDR    R1, b          5          MUL    R0, R0, R1
6          MUL    R0, R0, R1     6
7                                7
8   five   DEFW   5              8
9                                9
```

Literals can be any expression that the assembler can evaluate, for example, the following are all valid:

```
ADD R0, R0, #(1 + 2)
```

```
ADD R0, R0, #-2
```

```
ADD R0, R0, #(2 - 1)
```

Note that `MUL` cannot use literals. To get around this, first `MOV` the literal into a register and then use `MUL` with that register.

5.3 Status flags

ARM has some 1-bit status flags that are set after a `CMP` instruction as shown below:

Flag	Meaning
Negative	Previous result was negative
Zero	Previous result was zero
Carry	Previous add or subtract generated a carry
Overflow	The previous add or subtract overflowed and went out of range

You can get any data operation to alter the flags by appending `S` to the instruction. For example:

```
SUBS R0, R1, R2
```

The above command will subtract `R2` from `R1` and store the result in `R0`, but it will also set the flags according to the value of `R0`. For example, if the value in `R0` was negative after the operation, the `Negative` flag would be set.

5.4 Other useful arithmetic commands

The `RSB` command is reverse subtract, and is useful for negating a literal:

```
RSB R1, R0, #0; R1 = 0 - R0 = -R0
```

The `MLA` command is multiply and add. It can only use registers as operands.

```
MLA R1, R2, R3, R4 ; R1 = (R2 * R3) + R4
```

16

6   If and While

6.1   If statements in ARM

Given an if statement written in Java such as this:

```
1   if(condition)
2   {
3       actionStatement;
4   }
```

How would we convert that into ARM assembly? The best way is to do the following:

1. Check to see if the condition is else

2. If it is, then branch to the end of the if statement

3. Otherwise, perform the action statement

An example implementation may be:

```
1           CMP     R1, R2  ; Compare two registers
2           BGE     endif   ; The condition
3           MOV     R1, #0  ; Whatever the action
4                                   ; statement may be
5   endif
```

Note that it is often easier to implement the inverse of some comparisons. For example, if we were evaluating (a >= b) we could do either of these:

- 1. Check whether a is greater than b

  2. Check whether a is equal to b

  3. Work out how to branch away if both are false

- 1. Inverse the comparison and check if a is less than b

  2. Branch away if it is

6.2   If else statements

If else statements are much like if statements, except that if the condition is false then you branch to the else action, and if the condition is true, then you branch to after the else statement once you've executed the action. Here's an example:

Lets implement if(a==b) a += 1; else b += 1; :

```
1           LDR     R0, a
2           LDR     R1, b
3           CMP     R0, R1  ; Compare
4           BNE     else    ;If they aren't equal
5           ADD     R0, R0, #1
6           B       end
7   else    ADD     R1, R1, #1
8   end
```

6.3 While statements

Implementing a while statement in ARM assembler is very similar to implementing an if statement, except after you have executed the action statement, you branch back to the initial condition again, like so:

```
1           ; Compare two registers
2   start   CMP     R1, R2
3           ; The condition
4           BGE     endif
5           ; Whatever the action statement may be
6           ADD     R1, R1, #1
7           ; Branch back to the start
8           B       start
9   endif
```

In ARM, we can make *any* instruction have a conditional code on it. Consequently, we can get rid of a branch instruction, like so:

```
1           ; Compare two registers
2   start   CMP     R1, R2
3           ; Add only if the flag is less than
4           ADDLT   R1, R1, #1
5           ; Branch back to the start
6           BLT     start
7   endif
```

You can compare the result of an arithmetic instruction with zero by appending S to the instruction like so:

```
1   start   ADDS    R0, #1
2           BNE     start
3   endif
```

This will keep adding 1 to R0 until it reaches 0. This is often useful when translating for loops that count up to a number. Instead of counting up, you can count down using the SUBS instruction and only iterate if the result isn't zero.

Note that if the code in your while loop is being executed repeatedly, then it's probably worth optimising it!

7 Addresses and Addressing

When a processor references memory it needs to produce an address.

The address needs the same number of bits as the memory address. i.e. 32 in ARM

addressing modes - mechanisms for generating addresses

7.1 Direct Addressing

Direct addressing is a mode where the address is simply contained within the instruction.

This requires an instruction longer than the address size which is a problem because ARMs maximum bit length is 32.

So far, we assumed that direct addressing uses LDR/STR instructions, for example:

```
LDR    R0,    b
LDR    R1,    c
ADD    R0,    R0, R2
STR    R0,    a
```

This looks like direct addressing but on ARM it's 'faked' by the assembler as a pseudo-instruction

7.1.1 Problems with direct addressing

ARM: both instructions and addresses are 32 bits, but the instruction also specifies operation so it can't contain every possible address.

Solution: allow a register to contain an address, use the address in the register to do loads and stores.

This is **Register Indirect Addressing**

7.2 Register Indirect Addressing

The address is held in the register

It takes only a few bits to select a register (4 bits in the case of ARM R0-R15)

A register can (typically) hold an arbitrary address (32 bits in the case of ARM)

ARM has register indirect addressing

Example   loading a register from a memory location: LDR R0, b

Could be done using register indirect addressing:

```
ADR    R2,    b        Move the address of b into R2
LDR    R0,    [R2]     Use address in R2 to fetch the value of b
```

This is still a bit limited - addresses are:

- range limited (within ADR 'instruction')
- fixed

but:

- ADRL pseudo-op allows larger range (at a price)
- having addresses a variable once it is often used again
- variable are usually 'near' each other

7.2.1 Address Arithmetic

We can operate on registers, so we can:

- store/load/move addresses
- do arithmetic to calculate addresses

Rather than use e.g. extra ADD instructions, we often use **Base + Offset Addressing** - address addition done within the operand.

We have actually been using this all along: Base = PC register.

7.3 Offset Addressing

In offset addressing the address is calculated from a register value and a number.

The register specifier is just a few bits, The offset can be 'fairly small'.

With one register 'pointer' any of several variables in nearby addresses may be addressed.

ARM allows offsets of 12 bits in LDR/STR

These bits can be added or subtracted, for example:

```
LDR    R0,   [R1, #8]
STR    R3,   [R6, #-0x240]
LDR    R7,   [R2, #short-constant]
```

This provides a range of $\pm$ 4 kilobytes around a 'base' register

In practice this method is adequate for most purposes.

8 Strings, bit shifts, rotations and tables

8.1 Working with Strings

A String is just a list of bytes where each byte represents a character. The last byte in the list is the null byte which contains the value 0.

8.1.1 Loading a String

In order to load the String into a register, the ADRL command must be used, for example:

```
1   msg      DEFB    "My message", 0
2            ALIGN
3
4            ADRL    R0, msg
5            SVC     3
```

8.1.2 Finding the length of a String

In order to find the length of a String, we need only to loop over each byte in the String until we get to the null byte, and add up how many iterations we've done as we go along:

```
1            ADRL    R1, message
2            MOV     R2, #0
3   count    LDRB    R0, [R1, R2]
4            CMP     R0, #0
5            ADDNE   R2, R2, #1
6            BNE     count
7            STR     R2, length
```

Here, we use R1 to store the message, R0 to read each character, and R2 to keep the count.

8.1.3 Getting the index of a character in a String

In order to do the equivalent of
String.indexOf(character), we can do something similar to finding the length of the String, but we only need to loop until we get to the character we're looking for.

```
1            ADRL    R1, message
2            LDRB    R2, character
```

```
3           count   LDRB    R0, [R1] #1
4                   CMP     R0, #0
5                   BEQ     end
6                   CMP     R0, R2
7                   BNE     count
8           end     ADRL    R0, message
9                   SUB     R1, R1, R0
10                  SUB     R1, R1, #1
11                  STR     R0, index
```

8.2 Bit shifting and rotations

Shift operations move all the bits in a word in one direction or another. For example:

Shift direction	Resultant word
No shift	10011011
Left shift	00110110
Right shift	01001101

You might have noticed that a zero is appended to whatever side the bits are moving away from in order to ensure that the word is still the same number of bits as before.

Rotation operations are very similar to shift operations, except instead of padding the word with zeroes, the bit that was lost is appended to the other side of the word:

Rotation direction	Resultant word
No rotation	10011011
Left rotation	00110111
Right rotation	11001101

ARM has four different operations for shifting and rotations:

Mnemonic	Meaning	Function
LSL n	Logical shift left	Shifts left by n bits. Any bits that are lost are zeroed.
LSR n	Logical shift right	Shifts right by n bits. Any bits that are lost are zeroed.
ASR n	Arithmetic shift right	Shifts right by n bits. Any bits that are lost are set as the sign bit (to preserve the signed bit).
ROR n	Rotate Right	Rotates right by n bits.

8.3 Accessing a row in a table

We can use bit shifts to access a row in a table:

```
1               ADRL    R1, table
2               LDR     R2, [R1, R0, LSL #2]
3
4   table       DEFW    0
5               DEFW    3
6               DEFW    6
7               DEFW    9
8               DEFW    12
9               DEFW    15
```

Note how this snippet has an optimisation to use one less register during the looping process than the string length snippet does. It only uses one register to loop over the string and then works out the number of times it's looped after the character has been found.

Shifting to the left is like multiplying by two, and shifting to the right is like dividing by two.

21

```
10            DEFW    18
```

This snippet will allow us to access the row with an index stored in `R0`, so if `R0` was `2` then the output in `R2` would be `6`.

9   Stacks

Note, this section contains some known inaccuracies, including but not necessarily exclusively about the positioning of the stack pointer in relation to the stack. **Verify this information before you use it.**

A stack is a data structure that is used by ARM assembly (and in a variety of other computing applications) to help maintain the state of a program.

The only operations that can be performed on a stack are `push` and `pop`. The former adds an item onto the stack and the latter removes the last item from the stack. Because these are the only two operations, a Stack is a type of *Last In First Out* data structure.

The stack is usually initialised using the `DEFS` command to reserve a block of space in memory. This means that the stack will be a fixed size.

In ARM assembly, one register is reserved as the Stack Pointer (SP). This keeps track of the memory location that addresses the last item in the stack.

Stacks can be implemented where the stack pointer starts at the highest memory address in the range allocated to the stack, or where it starts at the bottom. In ARM assembly, it is usual to start the stack pointer at the highest address and work down.

9.1   Pushing to a stack

Lets initialise a stack of five memory locations:

The stack pointer will currently be pointed at `0x00000420`, so the next stackable memory location will be `0x00000416`.

If we want to *push* an item onto the stack, we need to do two things:

- Write the value to the current memory location addressed by the stack pointer.
- Decrement the stack pointer by `4`

22

The ARM command for a stack push is `PUSH`. An example might be:

`PUSH  #20`

Or if you wanted to push one or more registers:

`PUSH  R1, R2, R3`

If we executed the first of these commands (`PUSH #20`), the stack would look like this:

...	
20	⟵ 0x00000416 (SP)
0	⟵ 0x00000412
0	⟵ 0x00000408
0	⟵ 0x00000404
0	⟵ 0x00000400
...	

Executing `PUSH R2, R3` when the values of `R2` and `R3` were `523` and `8` respectively would result in a stack such as:

...	
20	⟵ 0x00000416
523	⟵ 0x00000412
8	⟵ 0x00000408 (SP)
0	⟵ 0x00000404
0	⟵ 0x00000400
...	

9.2   Poping from the stack

The ARM syntax to pop an item from the stack is very similar. To pop the last item from the stack into `R1` the command would be:

`POP  R1`

The stack would now look like:

```
              ...

              20          ← 0x00000416

              523         ← 0x00000412 (SP)

              8           ← 0x00000408

              0           ← 0x00000404

              0           ← 0x00000400

              ...
```

And the value of `R1` would be `8`. This works for multiple registers too just like the `PUSH` command does.

Note that we don't actually have to reset the memory location `0x00000408` since it will be overwritten on any subsequent push and since the stack pointer is behind it, it will never be read from.

9.3 Other ways of accessing the stack

It's important to remember that the stack is just a set of memory locations, and the stack pointer is just a normal register. This means that other commands for accessing memory can also be used to access and manipulate the stack. For example, take a look at this block:

```
1         LDR     R1, [SP]
2         ADD     SP, SP, #4
```

Of course, we could have used post indexed addressing in order to do the same thing: `LDR R1, [SP], #4`

This will load the value of the memory location addressed by the value inside the stack pointer into `R1` and then increment the stack pointer by 4. Essentially, this is doing the same thing as `POP R1` will do.

In order to do the same thing for pushing to the stack we can do:

```
1         SUB     SP, SP, #4
2         STR     R1, [SP]
```

Which is equal to `STR R1, [SP, #-4]!`

9.4 Stacks and method calls

It is important that when a program branches from one method to another, the values stored in the registers are not overwritten for when the program branches back to the original method.

This problem is resolved by pushing the values stored in the registers onto the stack before a branch is made, and then poping them off the stack once the method has finished executing. Alternately, the method that is being called can preserve the values in the registers by pushing and poping them to and from the stack at the start and the end of the method. Obviously, only one such method needs to be employed, so it's best to stick to one convention for the whole program.

10 Methods

The most basic way of calling a method in ARM assembly is by using the `BL` command. This will do two things:

1. Move the current value of the program counter into the link register.

2. Branch to the label defined in the instruction.

The method being called will use move the value of the link register back into the program counter once it has finished execution. Here's an example:

```
1
2          start   B main
3
4   add R1R2        ADD     R1, R1, R2
5          ; Move the value of the Link
6          ; Register back into the PC
7          MOV     PC, LR
8
9   main    MOV     R1, #4
10          MOV     R2, #2
11          BL      addR1R2
12
```

This program will move the values 4 and 2 into R1 and R2 respectively, and then branch to a method that will add the two registers together and store the result in R1.

10.1   Saving the value of registers

What if a method used lots of registers internally. We should make sure that the state of the registers is the same at the end of the method as they are at the start of the method, since otherwise there could be issues with corrupted data.

The easiest way to do this is to stack the parameters, like so:

```
1   method
2          ; First lets store the registers
3          ; we're going to use into memory
4          PUSH    {R0-R2}
5          ; ===========================
6          ; Do stuff with the registers
7          ; ===========================
8          POP     {R0-R2}
9          ; Branch back to the calling instruction
10         MOV     PC, LR
```

If a method calls another method, then the link register will need to be stacked too!

10.2   Passing parameters

Passing parameters to a method is easy; just add them to the stack! To do this, just load the variable into a register, then add it to the stack:

```
1          LDR     R0, myvar
2          STR     R0, [SP,#-4]!
```

When the method wants to access a parameter, it can just load the value in memory addressed by the stack pointer, adding an offset of #4 if the second parameter needs to be accessed, #8 for the third and so on. If the fourth parameter was to be loaded, then the instruction would be LDR R0 [SP, #12].

10.3 Stack frames

A stack frame is a set of memory locations on the stack that relate to one method call. They are often formatted in a way such as:

```
┌─────────────────────────┐
│           ...           │
├─────────────────────────┤
│    Method variables     │  ←── (SP)
├─────────────────────────┤
│      Link register      │
├─────────────────────────┤
│    Method parameters    │
├─────────────────────────┤
│           ...           │
└─────────────────────────┘
```

In order to access method variables from inside the method, you just do:

```
LDR R0, [SP, #offset*4]
```

To access method parameters, you do the same, using the offset of whatever method parameter you want to get.

To call a method, setting method parameters, you need to add them to the stack and then BL to the method:

```
LDR R0, param_1
STR R0, [SP, #-4]
...
LDR R0, param_n
STR R0, [SP, #n*-4]
BL methodName
```

To return to the parent method, remember to add the number of words used in the stack frame to the stack pointer. For example, if we used four parameters, didn't bother with stacking the Link Register and didn't stack any method variables, then we'd just need to remove four words (twenty bytes) from the stack pointer:

```
ADD SP, SP, #20
MOV PC, LR
```

11 Switch statements

Obviously, a switch statement could be compiled down to a series of `if else` statements and then converted into ARM assembly. However, if we had a long list of conditions, and our condition was either near the bottom of the list, or maybe even the `default` choice, then we'd need to test the switch expression against all of the conditions before being able to decide what to do.

This has the potential to be very inefficient, so thankfully, there is another way to do it. We can use the value of the switch variable to branch directly to the piece of code that we want to execute for that case.

This can be achieved using a table of values where the offset in the table corresponds to the next address to load into the PC (i.e. the branch location). For example:

```
1   ; Define a table of variables with
2   ; the locations to branch to
3   swtable DEFW    case0
4           DEFW    case1
5           DEFW    default
6           DEFW    case2
7   ; Load the base address of the table
8   ; into R1
9           ADR     R1, swtable
10  ; Load the value of R1 into the PC.
11  ; Offset the value in R1 by (R0*4)
12          LDR     PC, [R1, R0, LSL #2]
13
14  case0   ; Do stuff
15          B end
16
17  case1   ; Do stuff
18          B end
19
20  case2   ; Do stuff
21          B end
22
23  default ; Do stuff
24          B end
```

The LSL command is used to shift the bits in R1 left. The number of places shifted is equal to the value of R0. This has the effect of adding (4*R0) onto R1.

It's important to catch cases where the expression isn't in the table too, for example, if (following the previous example) the values we wanted to implement specific logic for were (0, 1, 3), then we would need to catch cases where the value was less than 0 and greater than 3.

```
1           CMP     R0, #0
2           BLT     default
3           CMP     R0, #3
4           BGT     default
```

The above would go before the table lookup would occur.

12   Types of values

12.1   Signed and unsigned integers

Lets assume that integers only have 8 bits, and so can hold 2^8 (256) different values. If we use a signed integer, the most significant bit would indicate whether the number was negative or not. This means that we loose a bit's worth of data and therefore can only represent the numbers $0 - 128$, though we do gain the ability to distinguish negative numbers from positive numbers.

As you should already know, using a signed integer is also known as 2's complement.

Note that the numbers $0 - 127$ are represented as the same as both signed and unsigned integers, however, when a number reaches the value 128, then as a signed integer it would mean -1 and as an unsigned integer it would mean 128.

It is important to note, that ARM has different types of compare instructions for signed and unsigned integers. We can use this to our advantage in the switch example outlined in the previous section (see page 26).

If we treat the register `R0` as a signed integer, we will have to check whether it is below `0` and whether it's above `4`. However, if we use unsigned comparisons, we only need to check whether the register is above `4` since any negative value will be represented as a positive value much larger than `4`!

```
1           CMP     R0, #4
2           BHI     default
```

12.2   Booleans in ARM assembly

Since a boolean value is either true or false, it is obvious that we only need a minimum of one bit to represent it. However, due to the difficulties of accessing just one bit in the ARM architecture, it actually make more sense to represent a boolean as a byte, where the least significant bit is either `1` or `0` depending on the value of the boolean.

This is easy, since it means that booleans are represented as the integers `0` and `1`. Testing a boolean value is easy too, since all you need to do is compare it to zero or one.

```
CMP  R1, #0
```

12.2.1   Logical operators in ARM

ARM assembly implements five logical operators:

- `AND` (logical AND)
- `ORR` (logical OR)
- `EOR` (logical XOR)
- `BIC` (bit clear)
- `MVN` (logical not)

All of these are self explanatory except from the bit clear one, which is explained using a table:

A	B	BIC A B
0	0	0
0	1	0
1	0	1
1	1	0

`BIC` is equivalent to `A & !B`. If B is high, then the output will be `0`, otherwise, it will be `A`.

13   Communicating with Peripherals

A peripheral is a device that is attached to a computer. This could be a keyboard, mouse, printer or arguably even the main memory of the computer.

There are two methods of interfacing with peripherals; polling and interrupts.

13.1   Polling

Polling is where the CPU will constantly check for an update from a peripheral. This involves using a loop to check the contents of a status register and then reading any data from the peripheral.

The status register is a memory mapped register that both the CPU and peripheral have access to. When the peripheral has data to transfer to the CPU, it will change the value of the status register to match a specific pattern (usually setting just one bit in the register). For example, the test pattern may be setting bit 7 high. Next time the CPU checks the status register, if bit 7 is high, then it will know there is data waiting for it.

If this is the case, then the CPU will then read the data register. This is also a memory mapped register, and will contain the data from the peripheral.

13.1.1 Memory mapping

A memory mapped register appears to be located at a specific memory location, but instead actually is situated inside the peripheral and is *mapped* so that it appears to be at a memory address.

13.1.2 Implementing polling

In this simple implementation of a polling loop, we're assuming that the `status_register` and `data_register` have already been set up as aliases to the correct memory locations.

```
1               ; Load the program relative
2               ; address into R1
3               ADR     R1, status_reg
4   loop
5               ; Use R1 as a pointer to
6               ; load the data in the status
7               ; register into R0
8               LDR     R0, [R1]
9               ; Test the content of the
10              ; status register
11              TST     R0, 0x80
12              ; Loop if the test failed
13              BEQ     loop
14              ; Load the content from the
15              ; data register into R0
16              ADR     R1, data_reg
17              LDR     R0, [R1]
```

13.1.3 The `TST` instruction

TST performs a bitwise `AND` on it's operands (a register and a literal test pattern). It then updates the flags for a subsequent conditional instruction. If the result of the AND was zero, then subsequent `EQ` conditional instructions will execute - i.e. it compares the result of the AND operation to zero.

Using the example in the polling implementation, lets execute the command `TST R0, 0x80`, first where the test condition is not met, and then when it is:

Content of R0	00000000
Test pattern	10000000
Result	00000000

Content of R0	10110100
Test pattern	10000000
Result	10000000

In the second run, you can see that the result isn't zero, so the test pattern must have been met.

13.2 Interrupts

Unlike the polling method, an interrupt is initiated by the peripheral rather than the CPU. It follows the following steps:

1. **Stop the program** This is often done by the interrupt automatically, and works just like a branch and link.

2. **Save important registers** The STMFD command can be used to save important register. This would include the current program status register (CSPR) and any registers modified by the interrupt handler.

3. **Run the interrupt handler** This would usually involve double checking the status register to make sure that the data is ready to be read and then reading the data inside the data register.

4. **Restore the saved registers** Use LDMFD to restore the registers just like we did to save them.

5. **Restart the program** Copy the Link register into the program counter and add on #4 so that the next instruction is executed and the program carries on as before.

13.2.1 An implementation of an interrupt

This is how an interrupt may be implemented:

```
1    ; Stage 1 has already happened
2    ; When we branched here
3    interrupt
4            ; Stage 2 - save registers
5            STMFD SP!, {R0, R1}
6            ; Stage 3 - run the interrupt
7            ; handler
8            ; Check the status_reg
9            ADR   R1, status_reg
10           LDRB  R0, [R1]
11           TST   R0, #0x80
12           ; Go to an error if the
13           ; TST failed
14           BEQ   error
15           ; Load the data_reg info
16           ADR   R1, data_reg
17           LDRB  R0, [R1]
18           ; Store the data_reg info
19           STR     R0, locationInMemory
20
21           ; Acknowledge the transfer
22           ADR   R1, acknowledge
23           MOV   R1, #1
24           STR     R0, [R1]
25           ; Step 4 - restore the
26           ; registers
```

```
27              LDMFD SP!, {R0, R1}
28              ; Step 5 - Branch back to
29              ; were we came from
30              SUBS  PC, LR, #4
```

13.2.2 Having multiple interrupts

An interrupt vector table can be used to select the correct address to branch to when multiple interrupts are in place. All that is needed is a table of pointers to various methods for handling the different peripherals. Then, it's easy to find the correct peripheral by using the value set in the status register as an index in the table.

13.2.3 How an `SVC` call works

We know that an `SVC` call is a software *interrupt*. This means there will be an interrupt handler that will be called to do something once the `SVC` instruction is executed. However, we include a parameter between 0 and 4 when we use a software interrupt, and this parameter defines the exact behaviour of the interrupt.

How then, does the interrupt handler know what parameter was used then? Well, these are the processes that go on when an `SVC` call happens:

1. Branch and link to the interrupt handler, save registers using `STMFD SP, R0-R2, LR`.

2. Load the link register into a register and take off `#4`. This will be the `SVC X` instruction. `LDR R0, [LR, #-4]`

3. Clear the top eight bits (the conditional code and the instruction code). Use `BIC R0, R0, 0xFF000000`

4. Logic shift right by 8 places (this will give us the 16 bit parameter, with the upper 16 bits all zeroed) `LSR R1, R0, #8`.

5. Load the start of the interrupt vector table into R2: `LDR R2, vector_table`

6. Now, we should load the interrupt into the PC with an offset of `R1 * #4` using `LDR PC, [R2, R1, LSL #4]`

7. We can finally restore the registers using `LDMFD SP, R0-R2, PC`. At the same time, we can set the PC back to where it was.

13.3 Direct Memory Access

Direct Memory Access (DMA) is a way of avoiding excess CPU overhead for data transfer between peripherals. If each interrupt takes 100 instructions to execute, and data is being sent byte by byte to the processor, then the process may use up a lot of CPU time, especially if there's lots of data being transferred.

The solution to this is to have an external, hardware based solution on the motherboard with access to memory. This will take care of the interrupts for us instead of the CPU. This allows peripherals to write directly to memory rather than having to go through the CPU. The hardware controller can give the CPU an interrupt once every kilobyte or so is written, which would reduce the load on the CPU by 1000 times.

DMA is good, since it reduces load on the CPU, and enables greater parallelization, however, it requires extra hardware which makes the system more complicated. Besides, the memory of the computer and the peripheral itself are still doing the same amount of work; only the load on the CPU has been lessened.

14 Operating systems

Most operating systems consist of four elements; a stack, a kernel, multiple libraries for programs to use and a user interface.

14.1 The Kernel

The kernel is responsible for abstracting the implementation details of the hardware that the operating system is sitting on top of away so that programs and applications can interface with it more easily. It also protects programs from each other, so that one program cannot access another programs memory space. This is an important security feature, and it is important that the kernel cannot be bypassed.

The kernel contains three *managers* for controlling different resources:

- **Process manager** - manages the CPU time between different processes to facilitate multitasking.

- **Memory manager** - assigns blocks of memory to different processes and ensures that each process can only access the memory that it tries to access.

- **Peripheral manager** - allows the CPU to interface with peripherals by providing device drivers for each kind of peripheral.

14.1.1 Kernel modes

Most kernels have multiple different processing *modes* that will allow programs different privileges. The default mode is usually the *user* mode where programs can only access memory in their allocated space. There is also a *privileged* mode where programs can access any memory address. ARM has an *interrupt* mode too handle interrupts from peripherals.

15 Methods of running code

15.1 Compilers and assemblers

The job of compilers and assemblers is to convert human readable source code into machine executable code (often referred to as *machine code*). There are a number of steps associated with compilation:

1. **Lexical analysis (Word)** is when the source code is broken down into tokens (groups of one or more characters that represent something, like a parameter, punctuation, operands etc). This is where comments are filtered out of the source code; they are not included in the tokens.

2. **Syntaxic analysis (Structure)** checks that the tokens are formed so that together, they form structures in the correct syntax for the language of the code. This means checking that method calls have the right number of parameters, semicolons are in the correct place, etc.

3. **Semantic analysis (Meaning)** makes sure that the code makes sense; it'll check the types of variables, check that programmer defined names and definitions are legal and so on.

4. **Code generation**. If the compilation process gets this far, then the code is all correct and will run. All that remains now is to convert it into machine readable code. This can

be in a variety of different formats, such as bytecode for Java or machine code if Assembly languages are being assembled.

5. **Optimisation (compiler step only)** is an important part of compilation. Only higher level languages have any optimisation built into their compilation process though; most assemblers don't do it (since they rely on the human programmer doing so).

15.2   Interpreters

An interpreter reads code and executes it directly. This code can be written either by hand in a language such as Python, or it can have been assembled or compiled such as with most assembly languages and compiled languages such as Java.

Here are some of the pros and cons about interpreters:

- Simple to implement.

- Code is very portable, the interpreter guarantees that certain functions will be available and they will be accessible in the same way on whatever system the interpreter is running on.

- Theoretically, interpreters are fairly secure, though exploits in the past have proven that implementing a secure interpreter isn't something to be taken lightly (take a look at Java!)

- Interpreters can be used to prototype new hardware, since they can emulate what the hardware would do.

- Debugging is very easy on an interpreter since a programmer would have full access to the stack and so could see the value of variables and memory locations all the time. It potentially also allows for code execution to be paused or stepped through.

- Interpreters are often seen as being slow since they are running on top of the *real hardware*, so there will invariably be some (often significant) overhead.

- It's often hard to simulate a real computer. While this may not always be desirable, if the interpreter is designed to mimic hardware exactly, then this is often very hard to implement.

16   Java Bytecode

When a `.java` file is compiled, the compiler (`javac`) produces a `.class` file. This is a file that contains java *bytecode*. This bytecode can then be run by a Java Virtual Machine (JVM). This is done by executing the instructions in the bytecode sequentially, *interpreting* each instruction as it comes. By doing this, the Virtual Machine can provide some guarantees of security, including the ability to execute code inside a sandbox if it is suspected to be malicious or untrustworthy.

One of the beautiful things about compiling code to an abstract assembly language such as bytecode is that in order to run it, all you need is a virtual machine. In this sense, all you need to run any Java code on any specific platform is a virtual machine specific to that platform. This makes languages that compile down to an abstract assembly language inherently portable and ubiquitously compatible.

16.1   Types of instruction

We are familiar with three address instructions (such as `ADD R1, R2, R3`) and two address instructions (`LDR R1, var`), but bytecode also brings one and even zero address instructions

into the mix. The creators of Java wanted to ensure that bytecode was as compact as possible so that it could be transferred over slow Internet connections as quickly as possible. Having zero or one address instructions was a way to shave off excess characters in the bytecode in order to minimise the bandwidth that bytecode would need to be sent over the Internet.

Instructions that take less than two operands will get their source and/or destination operands implicitly. In Java, they are the top one or two items on the stack. Code that is written in this way is often referred to as a *stack machine*.

For example, in order to add two numbers using a zero address instruction, the two numbers would have to be present on the stack, like so:

...
5
2	← SP
0
0
0
...

When the *ADD* command is executed, the two values would be popped, added together and then pushed again:

...
7	← SP
0
0
0
0
...

So, the `ADD` command in bytecode is really four ARM instructions; two pop's, one add and a push.

One address instructions are similar and are required in order to correctly populate the stack

before any processing can occur. `PUSH a` will push `a` onto the stack, and `POP a` will pop `a` off the stack. The bytecode to execute the expression $x = (a + b) * (c - d)$ would be:

```
1            PUSH a
2            PUSH b
3            ADD
4            PUSH c
5            PUSH d
6            SUB
7            MUL
8            POP x
```

16.2   Interpreting bytecode

It is possible to interpreter bytecode by having a large switch statement with all the different instructions in it and just matching the bytes against it.

The speed at which bytecode runs is sometimes a problem though. Due to the nature of one and zero address instructions, the stack is heavily used, which means that the processor wastes a lot of cycles merely waiting for memory addresses to be read. Luckily, due to the structure of the stack, the spatial locality of the memory addresses is high, and they are cached efficiently, however, execution can still be $10 - 100$ times slower than compiled C code.

16.2.1   Speeding up bytecode

It is possible to compile bytecode into native machine instructions. Java makes use of a technique called *Dynamic Class Loading* so that it only loads classes when they are needed (which minimises the overhead of starting up programs and makes the memory usage lighter). The Java Virtual Matching (JVM) will also compiles classes in real time using the Just In Time (JIT) compiler. This is initially slower than a simple interpreter since it takes time to compile classes before use, but they run far faster once compiled.

The JIT makes use of dynamic compilation in order to progressively enhance the speed of code as it runs more times. When a section of code is run the first time, the JIT will spend only the minimum amount of time compiling it so that it can be executed as soon as possible, however, if the code is run again, and again, then each time, the JIT will improve upon the compilation so that it runs faster. Since only around 10% of code in programs is run frequently, this is very efficient, since only that small percentage will be compiled in detail.

17   JVM memory usage

The memory usage of a running Java program can be broadly separated into three areas; the class area, the stack area and the heap area.

The class area is a static area of memory. It contains the method code and class variables (both primitive types and pointers to objects).

The stack area holds parameters that are passed between methods, any local variables used within the context of methods and return links.

The heap area is a section of memory where all the objects used in the program are stored in memory. All variables in the program that reference an object will point somewhere in the heap.

The memory usage of the JVM is different and can be broken down into three other areas; code, variable and stack.

17.1 Object instantiation

It's easy to create a variable in Java:

```
String myString;
```

This will create a pointer named `myString`. At the moment, since no object has been instantiated the pointer will have the value of `null`. However, we can use the Java reserved word `new` to instantiate a new object and point the variable at that:

```
myString = new String("I <3 COMP15111");
```

When this line is run, a new *String* object will be instantiated, and stored in memory locations inside the heap. A pointer to the String object will be placed in the *myString* variable.

As more objects are created during the running time of the program, the size of the heap grows and grows.

17.2 The format of the heap

A heap is similar to a stack in some ways. It is an array of consecutive memory locations, and their is a pointer, a *heap pointer* pointing to the next free memory location in the heap.

Each object in the heap is composed of three parts; a header that contains information about the object such as it's size, storage for the instance variables and a storage table of methods that the object contains.

17.3 Stale objects and garbage collection

When a variable is no longer needed, like an instance variable inside a method that has finished executing, or a variable that has been overwritten by another object, the actual object is still present in the stack, even though nothing is referencing it. These abandoned objects are referred to as *stale*.

In Java programs, the repeated process of instantiating a temporary object that subsequently becomes obsolete can result in the size of the heap growing very quickly to unmanageable levels. The space taken up by referenced objects in the heap may be small, but all the stale objects are taking up a lot of space. Garbage collection is a technique to remedy this. It contains a number of steps, outlined below:

1. First, program execution is stopped so that the variables used in the program and the heap don't change for the duration of the garbage collection.

2. The GC (garbage collector) then walks through the heap from top to bottom, and marks objects as *live* or *unreferenced*. This stage is called **marking** and can be very time consuming (especially if the heap is large).

3. The GC will then go along and delete all the unreferenced objects in the heap. This will remove all the stale objects, but will the heap spread out over more memory addresses than required (since there are gaps where the stale objects were), so the *size* of the heap hasn't been reduced.

4. The GC will then defragment the heap by moving the objects next to each other so that there are no gaps. References to the objects must be updated so that they point to the new locations of the objects in the heap.

This isn't the only way of implementing garbage collection. Other languages such as Python seek to avoid the 'stop-the-world' behaviour that the Java GC exhibits by using a technique called *reference counting* and reclaiming heap space as soon as objects are stale (and have no references). The Java GC does aim to reduce the time it spends doing its job though; it has different levels of garbage collection; major collections (that defragment and compact the heap) and minor

18 Implementing arrays

A primitive array is merely a set of succedent memory locations; each containing one item in the array. The variable containing the array will be a pointer to the first item in the array. Accessing a specific item in an array (termed *array indexing*) can be done in constant ($O(1)$) time, which is an important property of arrays. It's important to remember that the index starts from 0 and will need to be multiplied by the number of bytes in a word so that the correct memory location is requested for array indexing.

Here's some sample code for array indexing:

```
1            ; Implement 'element = array[index];'
2         LDR R0, index ; Load the index to get from
3                      the array
4         LDR R1, array ; Load the base address of
5                      the array
6         MOV R2, #4
7         MUL R2, R0, R2 ; Multiply the index by 4
8                       for 4 byte ARM words
9         LDR R3, [R1, R2] ; Load the requested array
10                          index
11        STR element, R3 ; Save the value in memory
```

Of course, we could save one instruction by using LSL instead of multiplying in a separate instruction. See section 8.2 for more info.

18.1 Checking bounds

You should check the bounds of an array before each array access (load or store). If the index is out of bounds, then you should branch to an error handling method.

This is sample code for loading an element at an index (`index`) of an array with a length of `length` and a base pointer of `basePointer`:

```
1         LDR R0, length
2         LDR R1, index
3         CMP R1, R0
4         BHI outOfBounds
5         LDR R2, basePointer
6         LDR R0, [R2, R1, LSL #2]
```

The HI condition code is used to treat the numbers in the registers as unsigned. This way, if the index is negative, then the MSB will be 1 and it will certainly be bigger than the `length`. See section 12.1 for more info.